Multimodal Fake News & Image Tampering Diagnostic Report

Generated on 2026-08-02 00:55:24 IST | Project Forensic Audit
Analysis performed at 2026-08-02 00:55:14 IST

1. Executive Summary & Verdict

Final Classification Verdict	**REAL**
Overall Multimodal Risk Score	0.3054 (0.0=Real, 1.0=Fake)
Prediction Confidence	77.53%
Processing Time	0.30 seconds

This submission was classified as likely authentic, with moderate confidence (77.5%). The submitted image contributed most to this verdict (57% of the fused score), primarily because simulated propagation graph shows a low proportion of suspicious/bot-like spreader accounts (9%). No image was submitted, so the image signal defaulted to neutral in the fusion.

2. Evidence Summary

• Simulated propagation graph shows a low proportion of suspicious/bot-like spreader accounts (9%).

3. Model Results & Fusion Weights

Modality	Model Engine	Risk Score	Weight	Relative Impact
Text Module	TF-IDF / BERT Transformer	0.1132	0.40	14.8%
Image Module	ResNet50 / ELA / pHash	0.5000	0.35	57.3%
Graph Module	PyG Graph Attention (GAT)	0.3403	0.25	27.9%

4. Analyzed Article Content

"ISRO has successfully completed another round of ground tests for critical systems of its Reusable Launch Vehicle program. Officials stated that the tests evaluated guidance, navigation, and thermal protection components under simulated launch conditions. The organization noted that the data collected will be used for future experimental flight mission"

5. Text Explanation — Top Contributing Tokens

SHAP/coefficient attribution identified the following tokens as most influential in the text-model verdict: **officials, noted**.

6. Propagation Analysis Summary

Simulated propagation graph: 45 nodes, 51 cascade edges, 8 detected communities, diffusion depth 6, 9% suspicious/bot-like accounts. Graph Attention Network propagation risk score: 0.3403.

7. Fusion Analysis

The final verdict is a weighted linear combination of the unimodal risk scores (see Section 3), calibrated via Platt scaling fit on the project's labeled demo corpus. Any modality not submitted for this analysis defaults to a neutral 0.5 score rather than being fabricated. The modality with the largest 'Relative Impact' in the table above contributed most to the final classification.

8. Final Decision & Recommendations

This content does not show strong indicators of fabrication based on the signals available to this system. As with any automated assessment, cross-verification against primary sources is still recommended before using this content for high-stakes decisions.