Multimodal Fake News & Image Tampering Diagnostic Report

Generated on 2026-08-02 00:58:09 IST | Project Forensic Audit
Analysis performed at 2026-08-02 00:57:58 IST

1. Executive Summary & Verdict

Final Classification Verdict	**FAKE**
Overall Multimodal Risk Score	0.5023 (0.0=Real, 1.0=Fake)
Prediction Confidence	50.37%
Processing Time	0.30 seconds

This submission was classified as likely fabricated or misleading, with limited confidence (50.4%). The article's text contributed most to this verdict (56% of the fused score), primarily because sensationalist/conspiratorial language detected: anonymous. Additional supporting evidence: text model attribution flagged these tokens as most fake-indicative: allegedly, anonymous, companies. No image was submitted, so the image signal defaulted to neutral in the fusion.

2. Evidence Summary

• Sensationalist/conspiratorial language detected: anonymous.
• Text model attribution flagged these tokens as most fake-indicative: allegedly, anonymous, companies.
• Simulated propagation graph shows a low proportion of suspicious/bot-like spreader accounts (9%).

3. Model Results & Fusion Weights

Modality	Model Engine	Risk Score	Weight	Relative Impact
Text Module	TF-IDF / BERT Transformer	0.6969	0.40	55.5%
Image Module	ResNet50 / ELA / pHash	0.5000	0.35	34.8%
Graph Module	PyG Graph Attention (GAT)	0.1942	0.25	9.7%

4. Analyzed Article Content

"A newly discovered Himalayan herbal tea has been scientifically proven to reverse aging within just seven days, according to anonymous researchers. The drink reportedly regenerates damaged organs, restores perfect eyesight, and permanently eliminates every known disease without side effects. Pharmaceutical companies are allegedly trying to suppress the discovery because it would end the global med..."

5. Text Explanation — Top Contributing Tokens

SHAP/coefficient attribution identified the following tokens as most influential in the text-model verdict: **allegedly, anonymous, researchers, according, industry, companies, effects, pharmaceutical**.

6. Propagation Analysis Summary

Simulated propagation graph: 45 nodes, 134 cascade edges, 5 detected communities, diffusion depth 5, 9% suspicious/bot-like accounts. Graph Attention Network propagation risk score: 0.1942.

7. Fusion Analysis

The final verdict is a weighted linear combination of the unimodal risk scores (see Section 3), calibrated via Platt scaling fit on the project's labeled demo corpus. Any modality not submitted for this analysis defaults to a neutral 0.5 score rather than being fabricated. The modality with the largest 'Relative Impact' in the table above contributed most to the final classification.

8. Final Decision & Recommendations

This content shows characteristics consistent with misinformation. Before sharing or acting on it: verify the claim against at least one established fact-checking organization or primary source, check whether the accompanying image appears elsewhere in a different context (reverse image search), and treat unattributed claims and emotionally-charged language as a caution signal.